

BDA - Assignment 7

Anonymous

Contents

Exercise 1. Linear model: drowning data with Stan	1
---	---

Exercise 2. Hierarchical model: factory data with Stan	5
--	---

```
library('aaltobda')
data(drowning)
library (rstan)
rstan_options (auto_write = TRUE)
options(mc.cores = parallel ::detectCores())
```

Exercise 1. Linear model: drowning data with Stan

1.

There were two mistakes in the given Stan code. The errors are colored in red and the correction is colored in blue

Firstly, this:

```
parameters {
  real alpha;
  real beta;
  real<upper=0>sigma;
}
```

should be written as:

```
parameters {
  real alpha;
  real beta;
  real<lower=0>sigma;
}
```

Secondly, this:

```
generated quantities {
  real ypred;
  ypred = normal_rng(alpha + beta*xsigma);
}
```

should be written as:

```
generated quantities {
  real ypred;
  ypred = normal_rng(alpha + beta*xpred*sigma);;
}
```

Below the full corrected Stan code

```

drowning = "
data {
    int<lower=0> N;           // number of data points
    vector[N] x;              // observation year
    vector[N] y;              // observation number of drowned
    real xpred;               // prediction year
}

parameters {
    real alpha;                // intercept
    real beta;                 // coefficients for predictors
    real<lower=0> sigma;       // error scale
}

transformed parameters {
    vector[N] mu;
    mu = alpha + beta*x;
}

model {
    y ~ normal(mu, sigma);    // likelihood
}

generated quantities {
    real ypred;
    ypred = normal_rng(alpha + beta*xpred, sigma);
}
"

```

2.

I determine tau for the weakly informative prior $N(0, \tau^2)$ by simulating 100000 rounds from normal distribution by changing the value of standard deviation so that 0.5% and 99.5% quantiles $\approx \pm 69$, Hence $Pr(-69 \leq beta \leq 69) = 0.99$

```

sim = rnorm(1000000, 0, 26.7)
print(quantile(sim, c(0.005, 0.995)))

##      0.5%    99.5%
## -68.52450  68.61133

```

Now, weakly informative prior $N(0, \tau^2)$ is determined by $N(0, 26.7^2)$

3.

The weakly informative prior, $N(0, 26.7^2)$, is implemented in *model* block of the below Stan code

```

library('aaltobda')
data(drowning)

stan_code =

```

```

data {
  int<lower=0> N;                      // number of data points
  vector[N] x;                          // observation year
  vector[N] y;                          // observation number of drowned
  real xpred;                           // prediction year
}

parameters {
  real alpha;                            // intercept
  real beta;                             // coefficients for predictors
  real<lower=0> sigma;                  // error scale
}

transformed parameters {
  vector[N] mu;
  mu = alpha + beta*x;
}

model {
  beta ~ normal(0, 26.7);    // weakly informative prior
  y ~ normal(mu, sigma);    // likelihood
}

generated quantities {
  real ypred;
  ypred = normal_rng(alpha + beta*xpred, sigma);
}
"
```

```
drowning_data = list(N=nrow(drowning),
                      x=drowning$year,
                      y=drowning$drownings,
                      xpred=2019)
```

```
fit_drowning = stan(model_code=stan_code, data=drowning_data, refresh=0)
```

```
## Warning: There were 1135 transitions after warmup that exceeded the maximum treedepth. Increase max_treedepth or, alternatively, set n_chains < 4.
```

```
## Warning: Examine the pairs() plot to diagnose sampling problems
```

```
monitor(fit_drowning)
```

```
## Inference for the input samples (4 chains: each with iter = 2000; warmup = 0):
##
##          Q5      Q50     Q95      Mean       SD   Rhat Bulk_ESS Tail_ESS
## alpha    414.8  1791.2 3035.9  1777.2  809.1      1     1067     1126
## beta     -1.4    -0.8   -0.1    -0.8    0.4      1     1067     1126
## sigma    21.4    25.7   32.2    26.1    3.3      1     1500     1124
## mu[1]   138.9   152.8  166.5   152.8    8.5      1     1337     1342
## mu[2]   138.7   152.0  165.1   152.0    8.1      1     1366     1342
## mu[3]   138.4   151.2  163.8   151.2    7.8      1     1399     1373
```

```

## mu[4]    138.0 150.4 162.4 150.4 7.5    1    1437    1465
## mu[5]    137.8 149.6 161.0 149.5 7.2    1    1482    1543
## mu[6]    137.5 148.8 159.7 148.7 6.8    1    1535    1494
## mu[7]    137.2 147.9 158.4 147.9 6.5    1    1600    1739
## mu[8]    137.0 147.1 157.1 147.1 6.2    1    1679    1770
## mu[9]    136.6 146.3 155.9 146.3 5.9    1    1774    1933
## mu[10]   136.1 145.5 154.7 145.4 5.7    1    1891    2149
## mu[11]   135.8 144.6 153.4 144.6 5.4    1    2034    2540
## mu[12]   135.3 143.8 152.2 143.8 5.2    1    2221    2551
## mu[13]   134.7 143.0 151.1 143.0 5.0    1    2458    2563
## mu[14]   134.2 142.2 150.0 142.2 4.8    1    2743    2537
## mu[15]   133.7 141.4 148.9 141.3 4.7    1    3078    2416
## mu[16]   133.1 140.6 147.9 140.5 4.5    1    3427    2465
## mu[17]   132.5 139.7 146.9 139.7 4.4    1    3684    2446
## mu[18]   131.6 138.9 146.1 138.9 4.4    1    3865    2544
## mu[19]   130.9 138.1 145.2 138.1 4.4    1    3927    2502
## mu[20]   130.1 137.2 144.5 137.2 4.4    1    3843    2528
## mu[21]   129.2 136.4 143.8 136.4 4.4    1    3645    2643
## mu[22]   128.1 135.5 143.1 135.6 4.5    1    3242    2535
## mu[23]   127.1 134.8 142.6 134.8 4.7    1    2826    2478
## mu[24]   126.1 134.0 142.0 134.0 4.8    1    2499    2511
## mu[25]   124.9 133.2 141.4 133.1 5.0    1    2261    2543
## mu[26]   123.7 132.4 140.9 132.3 5.2    1    2064    2445
## mu[27]   122.5 131.6 140.5 131.5 5.4    1    1904    2389
## mu[28]   121.3 130.8 140.0 130.7 5.7    1    1778    2287
## mu[29]   120.1 130.0 139.7 129.9 6.0    1    1677    2207
## mu[30]   118.8 129.1 139.2 129.0 6.2    1    1591    2009
## mu[31]   117.5 128.3 138.8 128.2 6.5    1    1521    1898
## mu[32]   116.2 127.5 138.5 127.4 6.8    1    1464    1775
## mu[33]   114.8 126.6 138.2 126.6 7.2    1    1415    1710
## mu[34]   113.5 125.7 138.0 125.8 7.5    1    1375    1665
## mu[35]   112.0 124.9 137.7 124.9 7.8    1    1334    1616
## mu[36]   110.7 124.0 137.5 124.1 8.1    1    1304    1622
## mu[37]   109.3 123.2 137.2 123.3 8.5    1    1279    1537
## ypred    74.5 121.5 166.8 121.3 28.1    1    3380    3482
## lp__   -137.8 -134.8 -133.7 -135.1 1.3    1    978    1060
##
## For each parameter, Bulk_ESS and Tail_ESS are crude measures of
## effective sample size for bulk and tail quantities respectively (an ESS > 100
## per chain is considered good), and Rhat is the potential scale reduction
## factor on rank normalized split chains (at convergence, Rhat <= 1.05).

```


```

par(mfrow=c(1,2))

draws = as.data.frame(fit_drowning)[,41]
hist(draws, breaks=40, pch = 19, lwd=0.1, cex=0.4, xlim=c(25,225),
     col='midnightblue', xlab='Posterior predictive histogram for year 2019')

draws = as.data.frame(fit_drowning)[,2]
hist(draws, breaks=40, pch = 19, lwd=0.1, cex=0.4, xlim=c(-2,0.5),
     col='midnightblue', xlab='beta')

```



Here we can see a almost exact match to the histograms provided in the assignment instructions. What is worth mentioning is that the given linear model with gaussian noise is not probably the best model here as the model output negative number of drownings.

Exercise 2. Hierarchical model: factory data with Stan

Separate model

In the separate model, each machine j will have its own mean μ_j and standard deviation σ_j .

Although the assignment asked us to use uniform priors (i.e. the default in Stan), there was a separate message on behalf of the course staff that recommended using weakly informative priors instead. Therefore, I use $N \sim (100, 20^2)$ as the weakly informative prior for the means and $\tau \sim Cauchy(0, 10^2)$ as the weakly informative prior for the standard deviations.

```
library('aaltobda')
data(factory)

factory_stan = "
data {
  int<lower=0> N;                      // number of data points
  int<lower=0> K;                      // number of groups
  int<lower=1,upper=K> x[N];           // group indicator
  vector[N] y;
}
parameters {
  vector[K] mu;                         // group means
  vector<lower=0>[K] sigma;             // group standard deviations
}
model {
  mu ~ normal(100, 20);                 // weakly informative prior for mean
  sigma ~ cauchy(0, 10);                // weakly informative prior for st.deviation
  y ~ normal(mu[x], sigma[x]);          // ypred for each model with their own means and st.deviations
}
generated quantities {
  real ypred;
  ypred = normal_rng(mu[6], sigma[6]); // ypred for sixth model with its own mean and st.deviation
}
```

```

factory_pooled = list(N = nrow(factory)*ncol(factory),
                      K = ncol(factory),
                      x = rep(1:6, nrow(factory)),      # [1,6] as a bracket
                      y = c(t(factory[,1:6])))

fit_sep = stan(model_code=factory_stan, data=factory_pooled, refresh=0)
monitor(fit_sep)

## Inference for the input samples (4 chains: each with iter = 2000; warmup = 0):
##
##          Q5   Q50   Q95  Mean   SD  Rhat Bulk_ESS Tail_ESS
## mu[1]    66.0  80.0  97.2  80.7  9.7    1     3063    1854
## mu[2]    95.4 105.7 115.4 105.6  6.2    1     2321    1560
## mu[3]    79.7  88.8  99.5  89.1  6.1    1     3967    2504
## mu[4]   103.8 111.1 117.6 111.0  4.5    1     2687    1732
## mu[5]    82.2  90.6  99.8  90.8  5.6    1     3107    1920
## mu[6]    75.1  88.2 103.5  88.7  9.0    1     3491    1901
## sigma[1] 12.4  20.0  39.3  22.3  9.4    1     2802    2013
## sigma[2]  7.7  12.3  23.8  13.6  5.5    1     3196    2070
## sigma[3]  8.1  13.2  25.3  14.5  6.0    1     3787    2194
## sigma[4]  5.0   8.4  16.5   9.3  3.9    1     3280    2474
## sigma[5]  6.9  11.2  21.6  12.4  5.0    1     3720    2255
## sigma[6] 12.1  19.1  36.7  21.1  8.8    1     3052    1967
## ypred    50.6  87.7 128.0  88.2 24.6    1     3451    3100
## lp__    -92.7 -87.0 -83.5 -87.4  2.8    1     1442    2119
##
## For each parameter, Bulk_ESS and Tail_ESS are crude measures of
## effective sample size for bulk and tail quantities respectively (an ESS > 100
## per chain is considered good), and Rhat is the potential scale reduction
## factor on rank normalized split chains (at convergence, Rhat <= 1.05).

```


Separate model i)

```

draws_sep_mu = as.data.frame(fit_sep)[,6]
hist(draws_sep_mu, breaks=100, pch = 19, lwd=0.1, cex=0.4,
     main='Separate model', col='midnightblue',
     xlab='The posterior distribution of the mean of the quality measurements of the sixth machine')
abline(v=mean(draws_sep_mu), col='firebrick', lwd=2.5)

```

Separate model




The posterior distribution of the mean of the quality measurements of the sixth machine

```
print(mean(draws_sep_mu))  
  
## [1] 88.65498  
  
#axis(side=1, at=seq(0,200, 25))
```

We can notice that the R-hat value obtained for mu[6] with the stan model is close to 1 so we can assume that chains have converged and indeed use the obtained model results to plot the posterior distribution of mu[6]. By looking at the distribution, we can observe that it is quite narrow and. The median is approximately 89.

Separate model ii)

```
draws_sep_pred = as.data.frame(fit_sep)[,13]  
hist(draws_sep_pred, breaks=100, pch = 19, lwd=0.1, cex=0.4,  
     main='Separate model', col='midnightblue',  
     xlab='The predictive distribution for another quality measurement of the sixth machine')  
abline(v=mean(draws_sep_pred), col='firebrick', lwd=2.5)
```



```
print(mean(draws_sep_pred))

## [1] 88.17024

#axis(side=1, at=seq(-200,2500,50))
```

The R-hat value of the generated quantity “ypred” is also close to 1 so we can assume that chains have converged and indeed use the obtained model results to plot the posterior distribution of “ypred”. Here we see that the predictive distribution for another quality measurement of the sixth machine is even narrower and has even higher kurtosis than above. The distribution of the prediction mass is quite concentrated around the range of 50-125. The median is approximately 89.

Separate model iii)

By using separate estimates for each of the quality measures, we assume that the experiments were performed independently and that each machine has its own model i.e. there is no link between the models. Without any quality measurement for the seventh machine, we cannot infer any posterior distribution of the mean of the quality measurements of this new machine.

Pooled model

For the pooled model, each machine j will have a common mean μ and standard deviation σ meaning that the parameter μ and σ are now single real valued parameters instead of vectors. However, I still use $N \sim (100, 20^2)$ as the weakly informative prior for the means and $\tau \sim Cauchy(0, 10^2)$ as the weakly informative prior for the standard deviations.

```
library('aaltobda')
data(factory)

factory_stan = "
```

```

data {
  int<lower=0> N;                                // number of data points
  vector[N] y;
}
parameters {
  real mu;                                         // Common mean
  real<lower=0> sigma;                            // Common standard deviation
}
model {
  mu ~ normal(100, 20);                          // weakly informative prior for mean
  sigma ~ cauchy(0, 10);                         // weakly informative prior for st.deviation
  y ~ normal(mu, sigma);                         // ypred for each model with their common means and st.deviations
}
generated quantities {
  real ypred;
  ypred = normal_rng(mu, sigma);    // ypred for sixth model with common mean and st.deviation
}
"

length = length(c(factory$V1, factory$V2, factory$V3, factory$V4, factory$V5, factory$V6))
factory_pooled = list(N=length,
                      y = c(factory$V1, factory$V2, factory$V3, factory$V4, factory$V5, factory$V6))

fit_pool = stan(model_code=factory_stan, data=factory_pooled, refresh=0)
monitor(fit_pool)

## Inference for the input samples (4 chains: each with iter = 2000; warmup = 0):
##
##          Q5      Q50     Q95   Mean     SD   Rhat Bulk_ESS Tail_ESS
## mu      87.6    93.0   98.5  93.0   3.2     1    2755    2561
## sigma   14.8    17.9   22.4  18.2   2.4     1    2899    2235
## ypred   62.5    93.1  123.4  93.3  18.4     1    3690    3891
## lp__ -102.7 -100.5 -99.9 -100.8  1.0     1    1835    2413
##
## For each parameter, Bulk_ESS and Tail_ESS are crude measures of
## effective sample size for bulk and tail quantities respectively (an ESS > 100
## per chain is considered good), and Rhat is the potential scale reduction
## factor on rank normalized split chains (at convergence, Rhat <= 1.05).


```

Pooled model i)

```

draws_pool_mu = as.data.frame(fit_pool)[,1]
hist(draws_pool_mu, breaks=100, pch = 19, lwd=0.1, cex=0.4,
     main='Pooled model', col='midnightblue',
     xlab='The posterior distribution of the mean of the quality measurements of the sixth machine')
abline(v=mean(draws_pool_mu), col='firebrick', lwd=2.5)

```



The posterior distribution of the mean of the quality measurements of the sixth machine


```
print(mean(draws_pool_mu))
```

```
## [1] 93.01198
```

For the pooled model, the obtained R-hat value for “mu” is also close to 1 and therefore the results obtained should be reliable. The posterior distribution of the mean of the quality measurements of the sixth machine resembles highly of normal distribution. The median is approximately 93.

Pooled ii)

```
draws_pool_pred = as.data.frame(fit_pool)[,3]
hist(draws_pool_pred, breaks=100, pch = 19, lwd=0.1, cex=0.4,
     main='Pooled model', col='midnightblue',
     xlab='The predictive distribution for another quality measurement of the sixth machine')
abline(v=mean(draws_pool_pred), col='firebrick', lwd=2.5)
```



The predictive distribution for another quality measurement of the sixth machine

```
print(mean(draws_pool_pred))
```

```
## [1] 93.25814
```

```
#axis(side=1, at=seq(0,200, 25))
```

Again, for the pooled model, the obtained R-hat value for “ypred” is also close to 1 and therefore ther results obtained should be reliable.

For the predictive distribution for another quality measurement of the sixth machine with the pooled model, we see again that the distribution is really normally distributed with some outliers. However, as we want to spesifically identify the predictive distribution for another quality measurement of the sixth machine, this model seems rather useless as it makes no disparity between different machines and measurements. The median is approximately 93.

Pooled iii)

By using the combined estimate averaging data from all quality measurement, we assume that the machines and measurements are quite similar and that the machines have no influence on quality measurements. Therefore, the prediction for the seventh machine simply returns global mean value, as the model again assumes the mean is shared among all machines. Subsuquently, the distribution of the mean of the quality measurements of the seventh machine is the same as the one obtained in question i) for the sixth machine.

Hierarchical model

As informed by the course staff, applying hyperprios is preferred over uniform prior. Hence, I will again use $N \sim (100, 20^2)$ as the weakly informative prior for the means and $\tau \sim Cauchy(0, 10^2)$ as the weakly informative prior for the standard deviations.

```

library('aaltobda')
data(factory)

factory_stan = "
data {
  int<lower=0> N;                      // number of data points
  int<lower=0> K;                      // number of groups
  int<lower=1,upper=K> x[N];           // group indicator
  vector[N] y;
}
parameters {
  real mu0;                            // prior mean
  real<lower=0> sigma0;                // prior std
  vector[K] mu;                      // group means
  real<lower=0> sigma;                 // common st.deviation
}
model {
  mu0 ~ normal(100,20);               // weakly informative prior for hierarchical mean (hyperprior)
  sigma0 ~ cauchy(0, 10);              // weakly informative prior for hierarchical sigma (hyperprior)
  mu ~ normal(mu0, sigma0);            // weakly informative prior for mean
  sigma ~ cauchy(0,10);                // weakly informative prior for st.deviation
  y ~ normal(mu[x], sigma);
}
generated quantities {
  real ypred;
  real mu7;
  // ypred for sixth model with individual means, common st.deviation and hyperprios
  ypred = normal_rng(mu[6], sigma);
  // posterior for the new seventh machine
  mu7 = normal_rng(mu0, sigma0);
}
"
factory_pooled = list(N = nrow(factory)*ncol(factory),
                      K = ncol(factory),
                      x = rep(1:6, nrow(factory)),      # [1,6] as a bracket
                      y = c(t(factory[,1:6])))

fit_hier = stan(model_code=factory_stan, data=factory_pooled, refresh=0)

## Warning: There were 8 divergent transitions after warmup. Increasing adapt_delta above 0.8 may help.
## http://mc-stan.org/misc/warnings.html#divergent-transitions-after-warmup

## Warning: Examine the pairs() plot to diagnose sampling problems

monitor(fit_hier)

## Inference for the input samples (4 chains: each with iter = 2000; warmup = 0):
##
##          Q5     Q50     Q95    Mean     SD   Rhat Bulk_ESS Tail_ESS
## mu0      84.1   93.2  102.5  93.3  5.6  1.00     2437    2545
## sigma0    4.7   11.0   22.3  12.0  5.7  1.01     1077     845

```

```

## mu[1]    70.6   81.3   92.5   81.3   6.6   1.00    2010    2359
## mu[2]    92.1  102.1  112.7  102.2   6.4   1.00    1935    2447
## mu[3]    79.9   89.5   99.0   89.5   5.8   1.00    3407    2590
## mu[4]    94.6  106.2  117.0  106.0   6.7   1.00    1764    1555
## mu[5]    81.7   91.0  100.6  91.1   5.9   1.00    3294    2657
## mu[6]    78.8   88.1   97.5   88.1   5.8   1.00    3078    2580
## sigma    11.8   14.6   19.2   15.0   2.3   1.00    2357    2396
## ypred    61.7   88.0  114.0  88.0  16.1   1.00    3579    3930
## mu7      70.0   93.2  116.4  93.2  14.3   1.00    3610    3378
## lp__   -114.8 -109.9 -107.1 -110.3   2.4   1.00    1200    1850
##
## For each parameter, Bulk_ESS and Tail_ESS are crude measures of
## effective sample size for bulk and tail quantities respectively (an ESS > 100
## per chain is considered good), and Rhat is the potential scale reduction
## factor on rank normalized split chains (at convergence, Rhat <= 1.05).


```

Hierarchical i)

```

draws_hier_mu = as.data.frame(fit_hier)[,8]
hist(draws_hier_mu, breaks=100, pch = 19, lwd=0.1, cex=0.4,
     main='Hierarchical model', col='midnightblue',
     xlab='The posterior distribution of the mean of the quality measurements of the sixth machine')
abline(v=mean(draws_hier_mu), col='firebrick', lwd=2.5)

```



The posterior distribution of the mean of the quality measurements of the sixth machine

```
print(mean(draws_hier_mu ))
```


```
## [1] 88.12872
```

For the hierarchical model, the obtained R-hat value for “mu[6]” is also close to 1 and therefore ther results obtained should be reliable.

For the posterior distribution of the mean of the quality measurements of the sixth machine with the hierarchical model we see that the distribution resembles a lot like the one with the pooled model. The difference is that the hierarchical model has fewer outliers and that the median is approximately 88, which is quite a bit less than in the pooled model and a lot closer to the non-pooled model. This can be explained by the fact that this model and the separate model take into account the specificities of the measurements of the sixth machine, that is, they account for machine specific means.

Hierarchical ii)

```
draws_hier_pred = as.data.frame(fit_hier)[,10]
hist(draws_hier_pred, breaks=100, pch = 19, lwd=0.1, cex=0.4,
     main='Hierarchical model', col='midnightblue',
     xlab='The predictive distribution for another quality measurement of the sixth machine')
abline(v=mean(draws_hier_pred), col='firebrick', lwd=2.5)
```



```
print(mean(draws_hier_pred))
```

```
## [1] 88.00372
```

Again, for the hierarchical model, the obtained R-hat value for “ypred” is also close to 1 and therefore the results obtained should be reliable.

For the predictive distribution for another quality measurement of the sixth machine with the hierarchical model we see that the distribution resembles a lot like the one with the pooled model. The difference is that the hierarchical model has fewer outliers and that it is more evenly distributed. The median is approximately 89, which again is less than in the pooled model and a lot closer to the non-pooled model.


Hierarchical iii)

```

draws_hier_7 = as.data.frame(fit_hier)[,11]
hist(draws_hier_7, breaks=100, pch = 19, lwd=0.1, cex=0.4,
     main='Hierarchical model', col='midnightblue',
     xlab='The posterior distribution of the mean of the quality measurements of the seventh machine.')
abline(v=mean(draws_hier_7), col='firebrick', lwd=2.5)

```

Hierarchical model



The posterior distribution of the mean of the quality measurements of the seventh machine

```
print(mean(draws_hier_7))
```

```
## [1] 93.20223
```

Again, for the hierarchical model, the obtained R-hat value for “mu7” is also close to 1 and therefore the results obtained should be reliable.

The mean is approximately 94, which is almost equal to the one in the pooled model. However, the spread is a lot narrower in comparison to the above distribution. This is because here we sample from the posterior of the hierarchical prior.