

BDA - Assignment 8

Anonymous

Contents

Exercise 1	1
Exercise 2	7
Exercise 3	10
Exercise 4	12
Exercise 5	12

```
library (rstan)
library(loo)
rstan_options (auto_write = TRUE)
options(mc.cores = parallel::detectCores())
```

Exercise 1

Separate model

Although the assignment asked us to use uniform priors (i.e. the default in Stan), there was a separate message on behalf of the course staff that recommended using weakly informative priors instead. Therefore, I use $N \sim (100, 20^2)$ as the weakly informative prior for the means and $\tau \sim Cauchy(0, 10^2)$ as the weakly informative prior for the standard deviations.

```
library('aaltobda')
data(factory)

factory_stan = "
data {
  int<lower=0> N;                      // number of data points
  int<lower=0> K;                      // number of groups
  int<lower=1,upper=K> x[N];           // group indicator
  vector[N] y;
}
parameters {
  vector[K] mu;                         // group means
  vector<lower=0>[K] sigma;             // group standard deviations
}
model {
  mu ~ normal(100, 20);                // weakly informative prior for mean
  sigma ~ cauchy(0, 10);                // weakly informative prior for st.deviation
  y ~ normal(mu[x], sigma[x]);         // ypred for each model with their own means and st.deviations
```

```

}

generated quantities {
  real ypred;
  vector[N] log_lik;
  ypred = normal_rng(mu[6], sigma[6]); // ypred for sixth model with its own mean and st.deviation
  for(i in 1:N) {
    log_lik[i] = normal_lpdf(y[i] | mu[x[i]], sigma[x[i]]); //log-likelihood
  }
}
"

factory_pooled = list(N = nrow(factory)*ncol(factory),
                      K = ncol(factory),
                      x = rep(1:6, nrow(factory)),      # [1,6] as a bracket
                      y = c(t(factory[,1:6])))

fit_sep = stan(model_code=factory_stan, data=factory_pooled, refresh=0)
monitor(fit_sep)

```

```

## Inference for the input samples (4 chains: each with iter = 2000; warmup = 0):
##
##          Q5   Q50   Q95  Mean   SD Rhat Bulk_ESS Tail_ESS
## mu[1]     67.1  80.3  97.3  81.0  9.0  1.00    3106    2156
## mu[2]     95.4 105.7 114.9 105.6  6.1  1.00    3026    1900
## mu[3]     79.6  88.8 100.2  89.2  6.4  1.00    2620    1522
## mu[4]    103.4 111.1 117.6 110.8  4.7  1.00    2288    1451
## mu[5]     81.8  90.7 100.4  90.8  5.8  1.00    2831    2033
## mu[6]     75.7  88.2 103.0  88.6  8.7  1.00    3011    1726
## sigma[1]    12.4  19.8  36.9  21.8  8.8  1.00    3270    2593
## sigma[2]     7.6  12.2  23.6  13.5  5.5  1.00    3061    2296
## sigma[3]     8.3  13.1  24.6  14.5  6.0  1.00    2964    1961
## sigma[4]     5.1   8.5  17.1   9.4  4.2  1.00    2170    1741
## sigma[5]     6.9  11.4  22.0  12.7  5.3  1.00    3234    2236
## sigma[6]     12.2  19.1  36.0  21.0  8.3  1.00    2876    1936
## ypred       51.5  87.8 127.8  88.5 24.0  1.00    3383    2896
## log_lik[1]   -4.6  -4.0  -3.5  -4.0  0.3  1.00    2355    2396
## log_lik[2]   -4.9  -4.0  -3.4  -4.0  0.5  1.00    4119    2946
## log_lik[3]   -5.0  -4.1  -3.5  -4.1  0.5  1.00    4171    3699
## log_lik[4]   -4.3  -3.5  -2.9  -3.5  0.4  1.00    2925    2338
## log_lik[5]   -5.0  -4.0  -3.4  -4.1  0.5  1.00    4282    3303
## log_lik[6]   -7.1  -5.3  -4.4  -5.5  0.9  1.00    4847    3356
## log_lik[7]   -4.9  -4.2  -3.8  -4.3  0.4  1.00    3847    2945
## log_lik[8]   -4.3  -3.5  -3.0  -3.6  0.4  1.00    2324    1881
## log_lik[9]   -4.3  -3.6  -3.2  -3.7  0.3  1.00    2197    1848
## log_lik[10]  -4.6  -3.6  -3.0  -3.7  0.5  1.00    3377    1920
## log_lik[11]  -4.4  -3.6  -3.1  -3.7  0.4  1.00    2772    2258
## log_lik[12]  -4.6  -4.0  -3.5  -4.0  0.3  1.00    2032    1443
## log_lik[13]  -4.9  -4.2  -3.8  -4.3  0.4  1.00    3847    2945
## log_lik[14]  -4.5  -3.8  -3.2  -3.8  0.4  1.00    3147    2437
## log_lik[15]  -4.3  -3.6  -3.1  -3.7  0.4  1.00    2134    1624
## log_lik[16]  -4.1  -3.3  -2.8  -3.4  0.4  1.00    2055    1313
## log_lik[17]  -5.3  -4.1  -3.5  -4.2  0.6  1.00    4942    3611
## log_lik[18]  -5.1  -4.3  -3.8  -4.4  0.4  1.00    4425    3221

```

```

## log_lik[19] -7.4 -5.5 -4.6 -5.7 0.9 1.00 6184 3641
## log_lik[20] -4.2 -3.5 -3.0 -3.6 0.4 1.01 2232 1884
## log_lik[21] -4.3 -3.6 -3.1 -3.6 0.4 1.00 2084 1533
## log_lik[22] -5.0 -3.8 -3.2 -3.9 0.6 1.00 4286 3128
## log_lik[23] -5.0 -4.0 -3.4 -4.1 0.5 1.00 4282 3303
## log_lik[24] -4.8 -4.1 -3.6 -4.2 0.4 1.00 2507 1512
## log_lik[25] -5.0 -4.2 -3.7 -4.2 0.4 1.00 3016 2578
## log_lik[26] -6.8 -4.7 -3.9 -4.9 0.9 1.00 3963 2807
## log_lik[27] -7.1 -5.0 -4.1 -5.2 1.0 1.00 4615 3165
## log_lik[28] -4.1 -3.3 -2.8 -3.4 0.4 1.00 2055 1313
## log_lik[29] -4.1 -3.5 -2.9 -3.5 0.4 1.00 2168 1928
## log_lik[30] -4.9 -4.2 -3.7 -4.2 0.4 1.00 2940 2303
## lp__ -92.7 -86.8 -83.5 -87.3 2.8 1.00 1360 1715
##
## For each parameter, Bulk_ESS and Tail_ESS are crude measures of
## effective sample size for bulk and tail quantities respectively (an ESS > 100
## per chain is considered good), and Rhat is the potential scale reduction
## factor on rank normalized split chains (at convergence, Rhat <= 1.05).

```

Pooled model

For the pooled model, each machine j will have a common mean μ and standard deviation σ meaning that the parameter μ and σ are now single real valued parameters instead of vectors. Once again I apply $N \sim (100, 20^2)$ as the weakly informative prior for the means and $\tau \sim \text{Cauchy}(0, 10^2)$ as the weakly informative prior for the standard deviations.

```

library('aaltobda')
data(factory)

factory_stan = "
data {
  int<lower=0> N; // number of data points
  vector[N] y;
}
parameters {
  real mu; // Common mean
  real<lower=0> sigma; // Common standard deviation
}
model {
  mu ~ normal(100, 20); // weakly informative prior for mean
  sigma ~ cauchy(0, 10); // weakly informative prior for st.deviation
  y ~ normal(mu, sigma); // ypred for each model with their common means and st.deviations
}
generated quantities {
  real ypred;
  vector[N] log_lik;
  ypred = normal_rng(mu, sigma); // ypred for sixth model with common mean and st.deviation
  for (i in 1:N) {
    log_lik[i] = normal_lpdf(y[i] | mu, sigma);
  }
}"

```

```

length = length(c(factory$V1, factory$V2, factory$V3, factory$V4, factory$V5, factory$V6))
factory_pooled = list(N=length,
                      y = c(factory$V1, factory$V2, factory$V3, factory$V4, factory$V5, factory$V6))

fit_pool = stan(model_code=factory_stan, data=factory_pooled, refresh=0)
monitor(fit_pool)

## Inference for the input samples (4 chains: each with iter = 2000; warmup = 0):
##
##          Q5     Q50    Q95   Mean    SD   Rhat Bulk_ESS Tail_ESS
## mu        87.7   93.2  98.6  93.2  3.3    1    2461    2069
## sigma      14.7   18.0  22.6  18.2  2.4    1    2170    2001
## ypred      62.2   92.5 123.6  92.7 18.8    1    3738    3704
## log_lik[1] -4.2  -4.0  -3.8  -4.0  0.1    1    2345    1893
## log_lik[2] -4.1  -3.8  -3.6  -3.8  0.1    1    1970    2015
## log_lik[3] -4.1  -3.8  -3.6  -3.8  0.1    1    1970    2015
## log_lik[4] -9.0  -7.3  -6.1  -7.3  0.9    1    2239    2203
## log_lik[5] -5.5  -4.9  -4.5  -4.9  0.3    1    2519    2365
## log_lik[6] -5.2  -4.7  -4.3  -4.7  0.3    1    2538    2141
## log_lik[7] -4.5  -4.2  -4.0  -4.2  0.2    1    2703    2316
## log_lik[8] -4.9  -4.5  -4.2  -4.5  0.2    1    2702    2226
## log_lik[9] -4.3  -4.0  -3.8  -4.0  0.1    1    2367    2030
## log_lik[10] -4.1  -3.9  -3.7  -3.9  0.1    1    2119    1825
## log_lik[11] -4.1  -3.9  -3.7  -3.9  0.1    1    2134    2064
## log_lik[12] -4.1  -3.8  -3.6  -3.8  0.1    1    1960    2018
## log_lik[13] -4.1  -3.8  -3.6  -3.8  0.1    1    1970    2015
## log_lik[14] -4.1  -3.9  -3.7  -3.9  0.1    1    2156    1802
## log_lik[15] -5.5  -4.9  -4.5  -4.9  0.3    1    2519    2365
## log_lik[16] -4.3  -4.0  -3.8  -4.1  0.1    1    2456    2048
## log_lik[17] -5.4  -4.9  -4.5  -4.9  0.3    1    2424    2092
## log_lik[18] -5.1  -4.6  -4.3  -4.7  0.3    1    2615    2087
## log_lik[19] -4.2  -3.9  -3.7  -4.0  0.1    1    2201    2050
## log_lik[20] -5.1  -4.6  -4.3  -4.7  0.3    1    2615    2087
## log_lik[21] -4.4  -4.1  -3.9  -4.1  0.2    1    2661    2077
## log_lik[22] -4.1  -3.8  -3.6  -3.9  0.1    1    1967    1998
## log_lik[23] -4.2  -4.0  -3.8  -4.0  0.1    1    2283    1944
## log_lik[24] -4.4  -4.1  -3.9  -4.1  0.2    1    2661    2077
## log_lik[25] -4.1  -3.8  -3.6  -3.8  0.1    1    1970    2015
## log_lik[26] -6.9  -5.8  -5.1  -5.9  0.5    1    2282    1902
## log_lik[27] -4.1  -3.8  -3.6  -3.8  0.1    1    1970    2015
## log_lik[28] -4.3  -4.0  -3.8  -4.0  0.1    1    2367    2030
## log_lik[29] -4.5  -4.2  -4.0  -4.2  0.2    1    2691    2254
## log_lik[30] -4.1  -3.9  -3.7  -3.9  0.1    1    2067    2142
## lp__     -102.9 -100.5 -99.9 -100.8 1.0    1    1612    1868
##
## For each parameter, Bulk_ESS and Tail_ESS are crude measures of
## effective sample size for bulk and tail quantities respectively (an ESS > 100
## per chain is considered good), and Rhat is the potential scale reduction
## factor on rank normalized split chains (at convergence, Rhat <= 1.05).
##
```

Hierarchical model

As informed by the course staff, applying hyperpriors is preferred over uniform priors. Therefore, I will one again use $N \sim (100, 20^2)$ as the weakly informative prior for the means and $\tau \sim Cauchy(0, 10^2)$ as the weakly informative prior for the standard deviations.

```
library('aaltobda')
data(factory)

factory_stan = "
data {
  int<lower=0> N;                      // number of data points
  int<lower=0> K;                      // number of groups
  int<lower=1,upper=K> x[N];           // group indicator
  vector[N] y;
}
parameters {
  real mu0;                            // prior mean
  real<lower=0> sigma0;                // prior std
  vector[K] mu;                       // group means
  real<lower=0> sigma;                 // common st.deviation
}
model {
  mu0 ~ normal(100,20);               // weakly informative prior for hierarchical mean (hyperprior)
  sigma0 ~ cauchy(0, 10);              // weakly informative prior for hierarchical sigma (hyperprior)
  mu ~ normal(mu0, sigma0);           // weakly informative prior for mean
  sigma ~ cauchy(0,10);                // weakly informative prior for st.deviation
  y ~ normal(mu[x], sigma);
}
generated quantities {
  real ypred;
  real mu7;
  vector[N] log_lik;
  // ypred for sixth model with individual means, common st.deviation and hyperprios
  ypred = normal_rng(mu[6], sigma);
  // posterior od the seventh machine with individual means, common st.deviation and hyperprios
  mu7 = normal_rng(mu0, sigma0);
  for (i in 1:N)
    log_lik[i] = normal_lpdf(y[i] | mu[x[i]], sigma); //log-likelihood
}

factory_pooled = list(N = nrow(factory)*ncol(factory),
                      K = ncol(factory),
                      x = rep(1:6, nrow(factory)),      # [1,6] as a bracket
                      y = c(t(factory[,1:6])))

fit_hier = stan(model_code=factory_stan, data=factory_pooled, refresh=0)

## Warning: There were 6 divergent transitions after warmup. Increasing adapt_delta above 0.8 may help.
## http://mc-stan.org/misc/warnings.html#divergent-transitions-after-warmup

## Warning: Examine the pairs() plot to diagnose sampling problems
```

```

monitor(fit_hier)

## Inference for the input samples (4 chains: each with iter = 2000; warmup = 0):
##
##          Q5     Q50    Q95   Mean    SD   Rhat Bulk_ESS Tail_ESS
## mu0      84.4   93.4  102.6  93.4   5.6    1     2488    2055
## sigma0    4.8    11.0   21.7  11.9   5.5    1     1052     800
## mu[1]     70.2   81.4   92.2  81.3   6.8    1     2014    1247
## mu[2]     92.8  102.5  112.4 102.4   6.0    1     2537    2751
## mu[3]     80.2   89.5   98.9  89.5   5.6    1     3125    2875
## mu[4]     95.2  106.2  117.3 106.2   6.8    1     2357    1919
## mu[5]     81.9   91.2  100.1  91.1   5.6    1     3051    2718
## mu[6]     78.5   88.3   98.1  88.3   6.0    1     2953    2359
## sigma     11.8   14.6   18.9  14.9   2.2    1     2417    2440
## ypred     62.0   88.3  115.1  88.4  16.3    1     3950    3862
## mu7       71.2   93.3  115.1  93.3  13.7    1     3193    3309
## log_lik[1] -4.1   -3.7   -3.4  -3.7   0.2    1     2450    2371
## log_lik[2] -4.9   -4.1   -3.6  -4.2   0.4    1     2820    3243
## log_lik[3] -4.7   -3.9   -3.6  -4.0   0.3    1     3684    3301
## log_lik[4] -4.1   -3.7   -3.4  -3.7   0.2    1     2156    2712
## log_lik[5] -4.6   -4.0   -3.6  -4.0   0.3    1     3119    3039
## log_lik[6] -7.8   -5.9   -4.7  -6.0   1.0    1     4158    2621
## log_lik[7] -4.8   -3.9   -3.6  -4.0   0.4    1     3675    3245
## log_lik[8] -4.2   -3.7   -3.5  -3.8   0.2    1     2205    2792
## log_lik[9] -4.0   -3.7   -3.4  -3.7   0.2    1     2513    2814
## log_lik[10] -4.9   -4.0   -3.5  -4.1   0.4    1     2210    3211
## log_lik[11] -4.2   -3.7   -3.5  -3.8   0.2    1     3074    2762
## log_lik[12] -4.1   -3.7   -3.4  -3.7   0.2    1     2762    3170
## log_lik[13] -4.8   -3.9   -3.6  -4.0   0.4    1     3675    3245
## log_lik[14] -4.6   -3.9   -3.5  -4.0   0.3    1     2591    3076
## log_lik[15] -4.0   -3.7   -3.4  -3.7   0.2    1     2312    2813
## log_lik[16] -4.6   -3.9   -3.5  -3.9   0.3    1     1952    2800
## log_lik[17] -4.7   -4.0   -3.6  -4.0   0.3    1     3506    3008
## log_lik[18] -5.2   -4.2   -3.7  -4.3   0.5    1     3593    3118
## log_lik[19] -8.8   -6.5   -4.9  -6.6   1.2    1     3441    1581
## log_lik[20] -4.0   -3.7   -3.4  -3.7   0.2    1     2004    2550
## log_lik[21] -4.1   -3.7   -3.4  -3.7   0.2    1     2092    2234
## log_lik[22] -4.2   -3.7   -3.5  -3.8   0.2    1     3089    2881
## log_lik[23] -4.6   -4.0   -3.6  -4.0   0.3    1     3119    3039
## log_lik[24] -4.6   -3.9   -3.5  -4.0   0.3    1     2967    2631
## log_lik[25] -5.0   -4.1   -3.6  -4.2   0.5    1     2209    1478
## log_lik[26] -5.2   -4.2   -3.7  -4.3   0.5    1     2966    3639
## log_lik[27] -6.0   -4.8   -4.0  -4.9   0.6    1     3840    3192
## log_lik[28] -4.6   -3.9   -3.5  -3.9   0.3    1     1952    2800
## log_lik[29] -4.0   -3.7   -3.4  -3.7   0.2    1     2320    2611
## log_lik[30] -4.7   -4.0   -3.6  -4.0   0.4    1     3911    3637
## lp__     -114.6 -109.9 -107.1 -110.3  2.3    1     1480    2595
##
## For each parameter, Bulk_ESS and Tail_ESS are crude measures of
## effective sample size for bulk and tail quantities respectively (an ESS > 100
## per chain is considered good), and Rhat is the potential scale reduction
## factor on rank normalized split chains (at convergence, Rhat <= 1.05).

```

Exercise 2

Note 1: As I use weakly informative prios, the resulting values for \hat{k} and \widehat{elpd}_{loo} can vary from the example solutions

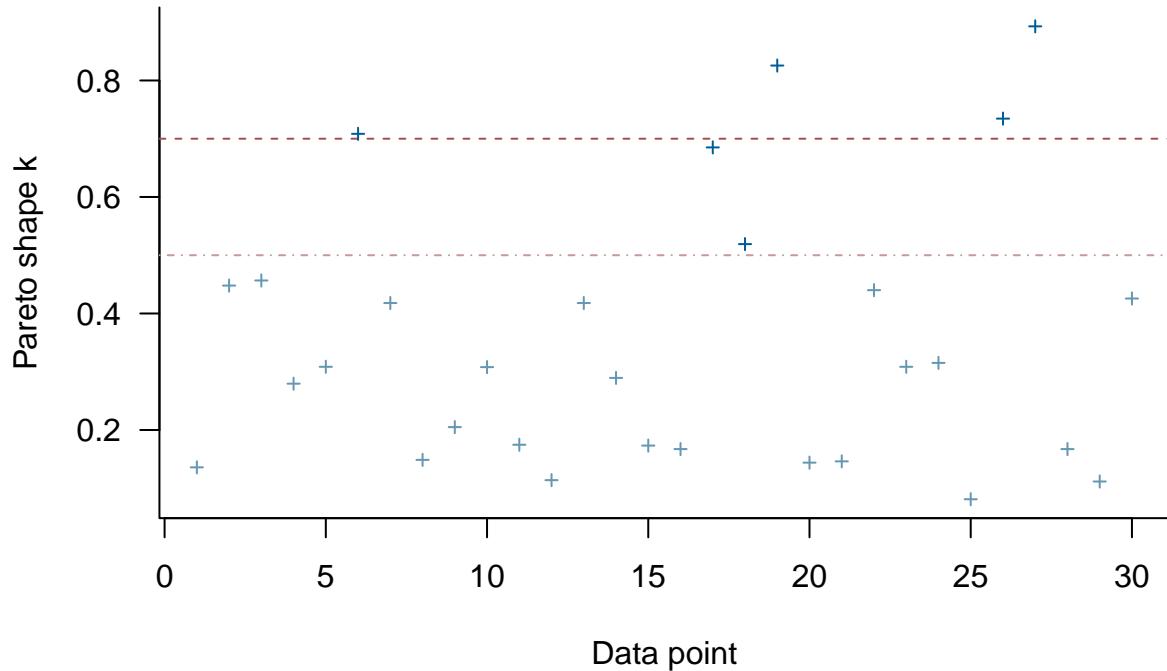
Note 2: To highlight the answers, I report some of the highlighted values as approximations because the knit to pdf conversion always changes the results a bit as the code chunks get re-run. The strictly correct values are, however, displayed in the print tables.

```
loo_sep = loo(fit_sep, cores=2)
print(loo_sep)
```

```
##
## Computed from 4000 by 30 log-likelihood matrix
##
##           Estimate   SE
## elpd_loo    -128.4 4.4
## p_loo        9.4 1.9
## looic      256.8 8.8
## -----
## Monte Carlo SE of elpd_loo is NA.
##
## Pareto k diagnostic values:
##                               Count Pct.  Min. n_eff
## (-Inf, 0.5]    (good)    24 80.0% 1417
## (0.5, 0.7]    (ok)      2  6.7% 992
## (0.7, 1]      (bad)     4 13.3%  57
## (1, Inf)     (very bad) 0  0.0% <NA>
## See help('pareto-k-diagnostic') for details.
```

```
plot(loo_sep)
```

PSIS diagnostic plot



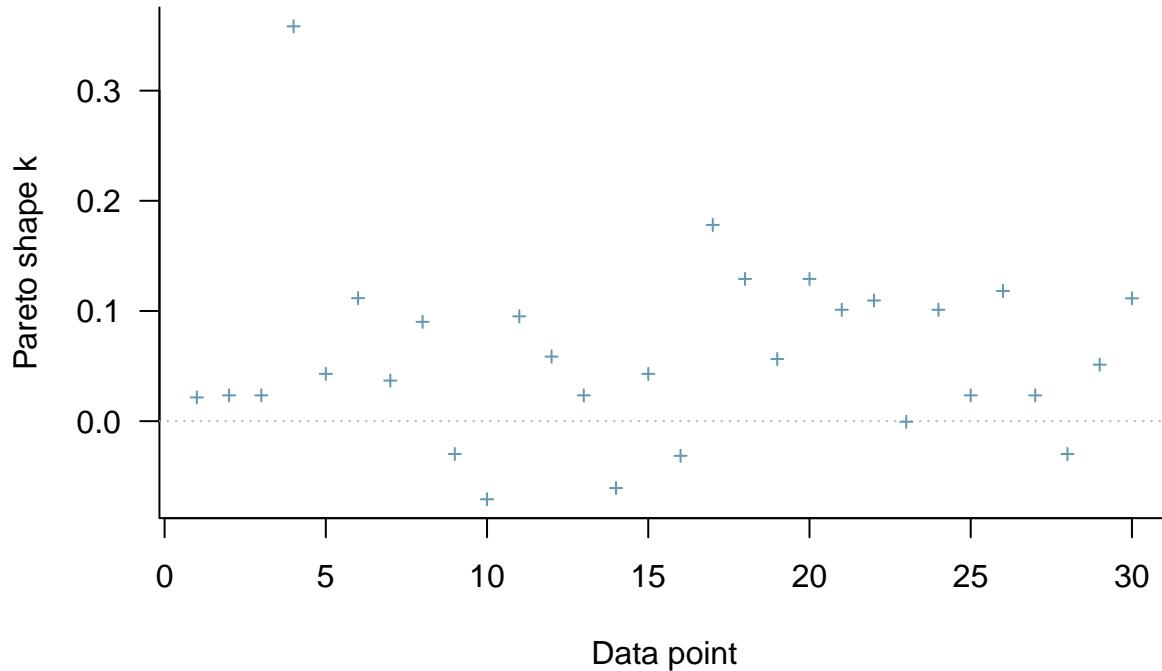
- PSIS – LOO \widehat{elpd}_{loo} value for the separate model ≈ -128.5
- For the separate model $\hat{k} > 0.7$

```
loo_pool = loo(fit_pool, cores=2)
print(loo_pool)
```

```
##
## Computed from 4000 by 30 log-likelihood matrix
##
##           Estimate   SE
## elpd_loo    -130.9 4.5
## p_loo        2.0  0.8
## looic       261.8 9.1
## -----
## Monte Carlo SE of elpd_loo is 0.0.
##
## All Pareto k estimates are good (k < 0.5).
## See help('pareto-k-diagnostic') for details.
```

```
plot(loo_pool)
```

PSIS diagnostic plot

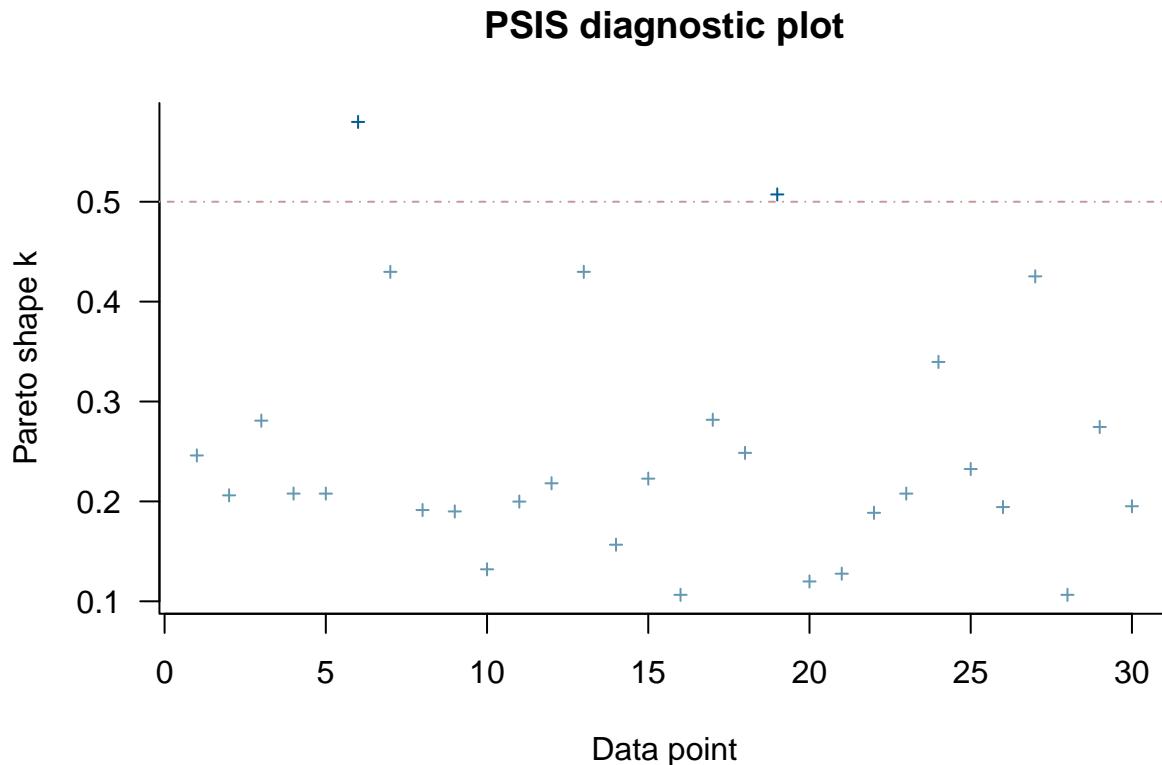


- $PSIS - LOO \widehat{elpd}_{loo}$ value for the pooled model ≈ -131.0
- For the pooled model $\hat{k} < 0.5$

```
loo_hier = loo(fit_hier, cores=2)
print(loo_hier)
```

```
##
## Computed from 4000 by 30 log-likelihood matrix
##
##           Estimate   SE
## elpd_loo    -126.9 4.6
## p_loo        5.7  1.7
## looic       253.8 9.2
## -----
## Monte Carlo SE of elpd_loo is 0.1.
##
## Pareto k diagnostic values:
##                               Count Pct. Min. n_eff
## (-Inf, 0.5]    (good)     28  93.3%  1600
## (0.5, 0.7]    (ok)       2   6.7%  332
## (0.7, 1]      (bad)      0   0.0% <NA>
## (1, Inf)     (very bad) 0   0.0% <NA>
##
## All Pareto k estimates are ok (k < 0.7).
## See help('pareto-k-diagnostic') for details.
```

```
plot(loo_hier)
```



- PSIS – LOO \widehat{elpd}_{loo} value for the hierarchical model ≈ -127.0
- For the hierarchical model $\hat{k} < 0.7$

Exercise 3

Compute Effective number of parameters for the separate model

Note: As I use weakly informative priors, the resulting values for \hat{p}_{loo} can vary from the example solutions

```
loo_sep = loo(fit_sep, r_eff=TRUE, cores=2)
print(loo_sep)
```

```
##
## Computed from 4000 by 30 log-likelihood matrix
##
##           Estimate   SE
## elpd_loo    -128.4 4.4
## p_loo        9.4 1.9
## looic      256.8 8.8
## -----
```

```

## Monte Carlo SE of elpd_loo is NA.
##
## Pareto k diagnostic values:
##                               Count Pct.    Min. n_eff
## (-Inf, 0.5]   (good)     24   80.0%  1417
## (0.5, 0.7]   (ok)       2    6.7%  992
## (0.7, 1]     (bad)      4   13.3%   57
## (1, Inf)     (very bad) 0    0.0% <NA>
## See help('pareto-k-diagnostic') for details.

```

- As we see from results, the \hat{p}_{loo} for the separated model is ≈ 9.6

```

loo_pool = loo(fit_pool, cores=2)
print(loo_pool)

```

```

##
## Computed from 4000 by 30 log-likelihood matrix
##
##           Estimate SE
## elpd_loo   -130.9 4.5
## p_loo        2.0  0.8
## looic      261.8 9.1
## -----
## Monte Carlo SE of elpd_loo is 0.0.
##
## All Pareto k estimates are good (k < 0.5).
## See help('pareto-k-diagnostic') for details.

```

- As we see from results, the \hat{p}_{loo} for the pooled model is ≈ 2.1

```

loo_hier = loo(fit_hier, cores=2)
print(loo_hier)

```

```

##
## Computed from 4000 by 30 log-likelihood matrix
##
##           Estimate SE
## elpd_loo   -126.9 4.6
## p_loo        5.7  1.7
## looic      253.8 9.2
## -----
## Monte Carlo SE of elpd_loo is 0.1.
##
## Pareto k diagnostic values:
##                               Count Pct.    Min. n_eff
## (-Inf, 0.5]   (good)     28   93.3%  1600
## (0.5, 0.7]   (ok)       2    6.7%  332
## (0.7, 1]     (bad)      0    0.0% <NA>
## (1, Inf)     (very bad) 0    0.0% <NA>
##
## All Pareto k estimates are ok (k < 0.7).
## See help('pareto-k-diagnostic') for details.

```

- As we see from results, the \hat{p}_{loo} for the hierarchical model is ≈ 5.9

Exercise 4

According to Vehtari et al (2016) The estimated shape parameter \hat{k} of the generalized Pareto distribution can be used to assess the reliability of the estimate:

1. If $\hat{k} < \frac{1}{2}$, the variance of the raw importance ratios is finite, and the estimate converges quickly.
2. If $\frac{1}{2} \leq \hat{k} < 1$, the variance of the raw importance ratios is infinite but the mean exists, and the convergence of the estimate is slower. If $\frac{1}{2} \leq \hat{k} < 0.7$ then we observe practically useful convergence rates and Monte Carlo error estimates, however, if $\hat{k} > 0.7$ we observe impractical convergence rates and unreliable Monte Carlo error estimates and the estimate may be biased (optimistic)
3. If $\hat{k} \geq 1$, the variance and the mean of the raw ratios distribution do not exist. The convergence rate is close to zero and the estimate is likely to be biased (optimistic)

As seen from the separate model results, the model is not reliable since $0.7 < \hat{k} \leq 1$ indicating that the convergence rate is close to zero and that the model is not reliable.

As seen from the pooled model results, all pareto \hat{k} estimates are good ($\hat{k} < 0.5$), indicating it converges quickly and that the model is reliable.

As seen from the hierarchical model results, all pareto \hat{k} estimates are good ($\hat{k} < 0.7$). In more detail, approximately 93% of the data points have \hat{k} values $\hat{k} < \frac{1}{2}$ and 7% have $\frac{1}{2} \leq \hat{k} < 0.7$. This indicates that the model converges fairly quickly and that the model is reliable.

(Visualizations for the \hat{k} values in section “Exercise 2”)

Exercise 5

We can observe from the results that the $PSIS - LOO \widehat{elpd}_{loo}$ values differ a bit between the three different models. According to these values, the hierarchical model has the highest $PSIS - LOO \widehat{elpd}_{loo}$ value at ≈ -127.0 , indicating that this model should be selected, as higher expected log pointwise predictive density corresponds to smaller predictive error

References

Vehtari, A., Gelman, A., and Gabry, J. (2017). Practical Bayesian model evaluation using leave-one-out cross-validation and WAIC. *Statistics and Computing*. 27(5), 1413–1432. doi:10.1007/s11222-016-9696-4. arXiv preprint: <http://arxiv.org/abs/1507.04544/>